

24

MPI – Point-to-point communication

- Communication between process 0 and process 1: [hello_1to0.c](#)

```
#include <stdio.h>
#include <string.h>
#include <sched.h> //with #define _GNU_SOURCE before
#include "mpi.h "

int main(int argc, char **argv)
{
    int          myRank, nbProcs, cpuId=-1, tag=50, nameLength;
    char          message[512], procName[MPI_MAX_PROCESSOR_NAME];
    MPI_Status    status;

    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myRank );
    MPI_Comm_size( MPI_COMM_WORLD, &nbProcs);

    MPI_Get_processor_name(procName, &nameLength);
    cpuId = sched_getcpu();
```

24

25

MPI – Point-to-point communication

- Communication between process 0 and others: [hello_1to0.c](#)

```
if ( myRank == 1 ) { // processes of rank 1
    sprintf( message, "Hello from %d on node %s-cpu%d!",
            myRank, procName, cpuId );
    MPI_Send( message, strlen(message)+1, MPI_CHAR,
            0, tag, MPI_COMM_WORLD ); }
else if ( myRank == 0 ) // process of rank 0
    MPI_Recv( message, 512, MPI_CHAR, 1,
            tag, MPI_COMM_WORLD, &status );
printf( "Proc %d: %s\n", myRank, message );
}

MPI_Finalize();
return 0;
}
```



25

26

MPI – Datatypes

- Similar to C, examples:

MPI datatype	C datatype
MPI_CHAR	char
MPI_INT	int
MPI_UNSIGNED	unsigned int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_BYTE	

- Complex data types

- ✧ MPI_PACKED performance of data communication
- ✧ Derived types: structure, column of matrix ...

26

27

MPI – Point-to-point communication

- Communication between process 0 and others: [hello_master.c](#)

```

#include <stdio.h>
#include <string.h>
#include <sched.h>
#include "mpi.h "
int main(int argc, char **argv)
{
    int myRank, nbProcs, cpuId=-1, src, tag=50, nameLength;
    char message[100], procName[MPI_MAX_PROCESSOR_NAME];
    MPI_Status status;
    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myRank );
    MPI_Comm_size( MPI_COMM_WORLD, &nbProcs);

    MPI_Get_processor_name(procName, &nameLength);
    cpuId = sched_getcpu();

```

27

28

MPI – Point-to-point communication

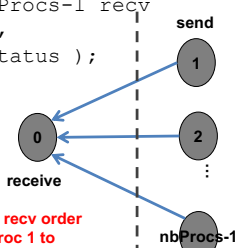
- Communication between process 0 and others: [hello_master.c](#)

```

if ( myRank != 0 ) { // processes of rank ≠ 0
    sprintf( message, "Hello from %d on node %s-cpu%d!",
            myRank, procName, cpuId );
    MPI_Send( message, strlen(message)+1, MPI_CHAR,
            0, tag+myRank, MPI_COMM_WORLD ); }
else // process of rank 0
    for ( src=1; src<nbProcs; src++ ) { // nbProcs-1 recv
        MPI_Recv( message, 100, MPI_CHAR, src,
            tag+src, MPI_COMM_WORLD, &status );
        printf( "%s\n", message );
    }

MPI_Finalize();
return 0;
}

```



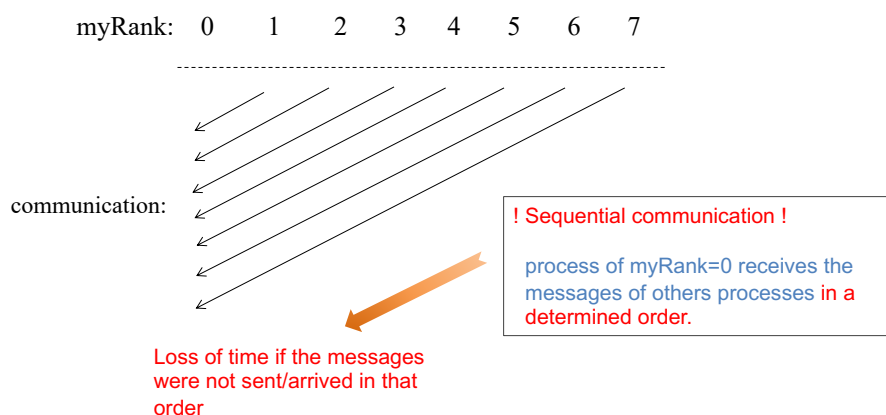
! Fixed recv order from proc 1 to nbProcs-1

28

29

MPI – Point-to-point communication

- Running of `hello_master.c` with 8 processes



29

30

MPI – Point-to-point communication

□ Modify the reception order of messages

```

if ( myRank != 0 ) {
    sprintf( message, "Hello from %d on %s-cpu%d!",
            myrank, procName, cpuId );
    MPI_Send( message, strlen(message)+1, MPI_CHAR,
            0, tag, MPI_COMM_WORLD ); }
else
    for ( src=1; src<nbProcs; src++ ) {
        MPI_Recv( message, 100, MPI_CHAR, MPI_ANY_SOURCE,
            tag, MPI_COMM_WORLD, &status );
        printf( "%s\n", message );
    }
MPI_Finalize();
return 0;
}

```

30

31

MPI – Point-to-point communication

□ Run the program with 8 processes

myRank: 0 1 2 3 4 5 6 7

communication:

! Sequential communication !

process of myRank=0 receives the messages of others processes in a undetermined order.
first arrived, first received


31

32

Point-to-Point Communication

- Communication between two processes
 - ✦ Source and Destination
- Message = header + data
- Data conversion if necessary
- Transmission mechanism

Source
Destination
Tag
Communicator
Data



32

33

Point-to-Point Communication

- Blocking communication
 - ✦ Blocking send and receive: `MPI_Send`, `MPI_Recv`
- Parameters of `MPI_Send` and `MPI_Recv`
 - ✦ Data address
 - ✦ Elements number of data
 - ✦ Type of data elements: `MPI_Datatype`
 - ✦ Source or Destination of the message (`MPI_ANY_SOURCE`)
 - ✦ Tag of message (`MPI_ANY_TAG`), may be used to indicate different type of message
 - ✦ Communicator (`MPI_COMM_WORLD`): `MPI_Comm`
 - ✦ Status: `MPI_Status` (`MPI_SOURCE`, `MPI_TAG`, `MPI_ERROR`)

33

34

MPI – Point-to-point communication

❑ Two side operation

➤ A Send must be matched by a Recv

- Safe program in blocking communication

```
MPI_Comm_rank (MPI_COMM_WORLD, myRank);
if (myRank == 0) {
    MPI_Send( sendBuf, count, MPI_INT, 1, tag1, MPI_COMM_WORLD );
    MPI_Recv( recvBuf, count, MPI_INT, 1,
              tag2, MPI_COMM_WORLD, &status );
}
else if (myRank == 1) {
    MPI_Recv( recvBuf, count, MPI_INT, 0,
              tag1, MPI_COMM_WORLD, &status );
    MPI_Send( sendBuf, count, MPI_INT, 0, tag2, MPI_COMM_WORLD );
}
```

34

35

MPI – Point-to-point communication

❑ Two side operation

➤ A Send must be matched by a Recv

- Safe program in blocking communication

```
MPI_Comm_rank (MPI_COMM_WORLD, myRank);
if (myRank == 0) {
    MPI_Send( sendBuf, count, MPI_INT, 1, tag1, MPI_COMM_WORLD );
    MPI_Recv( recvBuf, count, MPI_INT, 1,
              tag2, MPI_COMM_WORLD, &status );
}
else if (myRank == 1) {
    MPI_Send( sendBuf, count, MPI_INT, 0, tag2, MPI_COMM_WORLD );
    MPI_Recv( recvBuf, count, MPI_INT, 0,
              tag1, MPI_COMM_WORLD, &status );
}
```

➤ possible dead lock if we exchange the order of Recv and Send ! (depending on the implementation of MPI)

35

36

MPI – Blocking communication

□ Features

- Completion of `MPI_Send` means send variable can be reused
- Completion of `MPI_Recv` mean receive variable can be read
- Cause synchronization -> Increase communication time
- Affect the performance of parallel program

□ Solution (if there are many communication between two processes)

- Non-blocking communication

36

37

MPI – Non-blocking communication

□ Operation in 2 steps

- Request: `MPI_Isend`, `MPI_Irecv`

```

MPI_Isend(buf, count, datatype, dest,
          tag, comm, &request);
MPI_Irecv(buf, count, datatype, src,
          tag, comm, &request);

```
- Completion: `MPI_Wait(&request, &status);`
- Test of completion:


```

MPI_Test(&request, &flag, &status);

```
- Avoid dead lock
- Allow communication / computation overlapping
- Persistent request can be used if many communication

37

38

MPI – Non-blocking communication

□ Example

```
#include <stdio.h>
#include <string.h>
#include <mpi.h>

int main( int argc, char **argv ) {
    int myRank=-1, nbProcs=0;
    char sendMsg[128], recvMsg[128];

    MPI_Request requestS, requestR;
    MPI_Status status;
    use with nbProcs = 2!

    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myRank );
    MPI_Comm_size( MPI_COMM_WORLD, &nbProcs );

    sprintf(sendMsg, "Hello from proc %d/%d!", myRank, nbProcs);
```

38

39

MPI – Non-blocking communication

□ Example

```
if (myRank == 0) {
    MPI_Isend(sendMsg, strlen(sendMsg)+1, MPI_CHAR, 1,
        100, MPI_COMM_WORLD, &requestS);
    MPI_Irecv(recvMsg, 128, MPI_INT, 1,
        200, MPI_COMM_WORLD, &requestR);
}
if (myRank == 1) {
    MPI_Isend(sendMsg, strlen(sendMsg)+1, MPI_CHAR, 0,
        200, MPI_COMM_WORLD, &requestS);
    MPI_Irecv(recvMsg, 128, MPI_CHAR, 0,
        100, MPI_COMM_WORLD, &requestR);
    nbProcs = 2!
}
MPI_Wait(&requestR, &status);
printf("Proc %d: received message: %s\n", myRank, recvMsg);
MPI_Finalize();
return 0;
}
```

39

40

MPI – Non-blocking communication

□ Example

```
if (myRank == 0) {
    MPI_Isend(sendMsg, strlen(sendMsg)+1, MPI_CHAR, 1,
              100, MPI_COMM_WORLD, &requestS);
    MPI_Recv(recvMsg, 128, MPI_INT, 1,
              200, MPI_COMM_WORLD, &status);
}

if (myRank == 1) {
    MPI_Isend(sendMsg, strlen(sendMsg)+1, MPI_CHAR, 0,
              200, MPI_COMM_WORLD, &request1);
    MPI_Recv(recvMsg, 128, MPI_CHAR, 0,
              100, MPI_COMM_WORLD, &status);
}
```

mixte MPI_Isend and
MPI_Recv is possible.

40

41

41

42

MPI – Point-to-Point Communication

□ Exercise: Computing of PI – Numerical integration principle

➤ Mathematical formula

$$\pi = \int_0^1 \frac{4}{1+x^2} dx$$

➤ Rectangle rule

☆ Let n be the number of rectangles


☆ Let h be the width of rectangles

☆ We have:

$$h = \frac{1}{n}$$

➤ An approximation of π :

$$\sum_{i=0}^{n-1} h \frac{4}{1 + (ih + 0,5h)^2}$$



42

43

MPI – Point-to-Point Communication

□ Exercise: Sequential computing of PI

```
#include <stdio.h>
int main(int argc, char **argv)
{
    int i, nbRects=1000000; // read from keyboard in program
    double x, h, sum=0.0, pi=0.0;

    printf("Please input the number of rectangles of [0-1]: ");
    scanf("%d", &nbRects); // if (argc>1) nbRects = atoi(argv[1]);

    h = 1.0 / nbRects ;
    for (i=0; i<nbRects; i++) {
        x = (i+0.5)*h;    sum += 4.0 / (1.0 + x*x);
    }

    pi = h * sum ;
    printf("Pi is approximatly: %.16f\n", pi) ;
    return 0;
}
```

double PI25DT =
3.141592653589793238462643;

43

44

MPI – Point-to-Point Communication

□ Exercise: Parallel computing of PI – Send/Recv

```
#include <stdio.h>
#include <mpi.h>

int main(int argc, char **argv)
{
    int    i, nbRects= 1000000, rectsPerProc=0;
    int    myDeb=0, myEnd=0;
    double x, h, mySum=0.0, pi=0.0;

    int     myRank, nbProcs, tag=50;
    MPI_Status status;

    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myRank );
    MPI_Comm_size( MPI_COMM_WORLD, &nbProcs);
```

44

45

MPI – Point-to-Point Communication

□ Exercise: Parallel computing of PI – Send/Recv

```
if (myRank==0) { // process of rank 0 does inputs
    scanf("%d%c", &nbRects);

    rectsPerProc = nbRects / nbProcs;

    // rank 0 sends rects_per_proc to the other
    for (i=1; i<nbProcs; i++)
        MPI_Send(&rectsPerProc,1,MPI_INT,i,tag,MPI_COMM_WORLD);
}
else // process of rank#0 receives rects_per_proc from rank 0
    MPI_Recv(&rectsPerProc, 1, MPI_INT, 0, tag,
            MPI_COMM_WORLD, &status);

h = 1.0 / (rectsPerProc*nbProcs);
```

45

46

MPI – Point-to-Point Communication

❑ Exercise: Parallel computing of PI – Send/Recv

```

myDeb=myRank*rectsPerProc;  myEnd=myDeb+rectsPerProc;
for (i=myDeb; i<myEnd; i++) {
    x = (i+0.5)*h;    mySum += 4.0 / (1.0 + x*x);
}

if (myRank != 0)
    MPI_Send(&mySum, 1, MPI_DOUBLE, 0, tag, MPI_COMM_WORLD);
else { // rank 0
    for (i=1; i<nbProcs; i++) {
        MPI_Recv(&pi, 1, MPI_DOUBLE, MPI_ANY_SOURCE,
                tag, MPI_COMM_WORLD, &status);
        mySum = mySum + pi;
    }
    pi = mySum * h; printf("Pi is approximatly %0.16f\n", pi);
}
MPI_Finalize();  return 0;
}

```

46

47

Exercise: Calculation of PI – Send/Recv

❑ Resume

- The computation of the `sum` is well distributed.
- **Sequential communication** involves all processes at the beginning and the end of program.
 - ➔ May be improved by using **collective communication**

- ❑ To show more collective communication functions, we assume that only process of `rank 0` has the value of `nbRects` at the beginning of the program

47

48

48

49


MPI – Point-to-Point Communication

Exercise - Dot product

- Compute $d = X \cdot Y$, $X, Y \in \mathbb{R}^n$, n big number

$$d = \sum_{i=0}^{n-1} X_i Y_i, i = 0, \dots, n-1; X_i, Y_i \in \mathbb{R}$$

- Parallel algorithm - Example with 4 processes



49

50

50

51

Collective Communication

- What is it?
 - ✧ Communication involving all processes of a group
- Objective
 - ✧ Increase the performance of parallel program
- How?
 - ✧ By reduce of idle processes \implies decrease the communication time
- Use cases
 - ✧ When I/O
 - ✧ Parallel algorithms need collective communication


51

52

MPI – Collective communication

□ Broadcast (diffusion)

- A process (`root`) has a message to send to others



```
int MPI_Bcast(void *msg, int count,
             MPI_Datatype datatype, int root, MPI_Comm comm);
```


```
int send_buf[4];
//root initialize the send_buf
MPI_Bcast(send_buf, 4, MPI_INT, 0, MPI_COMM_WORLD);
```

52

53

MPI – Collective communication

□ Broadcast: Possible implementation



53

54

MPI – Collective communication

□ Broadcast - Broadcast of the dimension of a image

```
int myRank, nbProcs, dims[2], i, tag=30;
...
MPI_Comm_rank( MPI_COMM_WORLD, &myRank );
MPI_Comm_size( MPI_COMM_WORLD, &nbProcs );
if (myRank == 0) {
    /* Fill dims */
    for ( i=1; i<nbProcs; i++)
        MPI_Send( dims, 2, MPI_INT, i, tag, MPI_COMM_WORLD);
}
else
    MPI_Recv(dims, 2, MPI_INT, 0, tag, MPI_COMM_WORLD);
```

Point-to-point communication:
Number of steps - $O(nbProcs)$

```
int myRank, nbProcs, dims[2], i, tag=30;
...
MPI_Comm_rank( MPI_COMM_WORLD, &myRank );
MPI_Comm_size( MPI_COMM_WORLD, &nbProcs );
MPI_Bcast(dims, 2, MPI_INT, 0, MPI_COMM_WORLD);
```


Collective communication:
Number of steps - $O(\log_2(nbProcs))$
with binary tree

54

55

MPI – Collective communication

□ Scatter – Data distribution




```
int MPI_Scatter( void *sendbuf, int sendcnt, MPI_Datatype sendtype,
                void *recvbuf, int recvcnt, MPI_Datatype recvtype,
                int root, MPI_Comm comm );
```

55

56

MPI – Collective communication

□ Gather – Data fusion (fusion)




```
int MPI_Gather( void *sendbuf, int sendcnt, MPI_Datatype sendtype,
               void *recvbuf, int recvcnt, MPI_Datatype recvtype,
               int root, MPI_Comm comm );
```

56

57

MPI – Collective communication

□ Gather: Possible implementation




57

58

MPI – Collective communication

Reduce – Gather + operation on data (réduction)




```
int MPI_Reduce( void *operand, void *result,
               int count, MPI_Datatype datatype, MPI_Op op,
               int root, MPI_Comm comm );
```

58

59

MPI – Collective communication

Reduce: Possible implementation



59

60

Collective Communication

□ Reduce operations

Operation	Meaning	Operation	Meaning
MPI_MAX	Maximum	MPI_LOR	logical OR
MPI_MIN	Minimum	MPI_BOR	bitwise OR
MPI_SUM	Sum	MPI_LXOR	XOR
MPI_PROD	Product	MPI_BXOR	
MPI_LAND	Logical AND	MPI_MAXLOC	max or min +
MPI_BAND	bitwise AND	MPI_MINLOC	index

60

61

MPI – Collective communication

□ Exercise: Parallel computing of PI

```

#include <stdio.h>
#include <mpi.h>

int main(int argc, char **argv)
{
    int    i, nbRects= 1000000, rectsPerProc=0;
    int    myDeb=0, myEnd=0;
    double x, h, mySum=0.0, pi=0.0;

    int     myRank, nbProcs, tag=50;
    MPI_Status status;

    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myRank );
    MPI_Comm_size( MPI_COMM_WORLD, &nbProcs);

```

61

62

MPI – Collective communication

□ Exercise: Parallel computing of PI

```

if (myRank==0) { // process of rank 0 does inputs
    printf("Please input the number of rectangles for [0, 1]: ");
    scanf("%d%c", &nbRects);

    rectsPerProc = nbRects / nbProcs;
}

MPI_Bcast(&rectsPerProc, 1, MPI_INT, 0, MPI_COMM_WORLD);

h = 1.0 / (rectsPerProc*nbProcs);

```

62

63

MPI – Collective communication

□ Exercise: Parallel computing of PI

```

myDeb = myRank * rectsPerProc;
myEnd = my_deb + rectsPerProc;

for (i=myDeb; i<myEnd; i++) {
    x = (i+0.5) * h;
    mySum += 4.0 / (1.0 + x*x);
}
mySum = h * mySum ;

MPI_Reduce(&mySum, &pi, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);

if (myRank==0)
    printf("Pi is approximatly %0.16f\n", pi);

MPI_Finalize();
return 0;
}

```

63

64

MPI – Collective communication

- ❑ Barrier synchronization
 - Make a appointment for all processes
- ❑ Use case: time measurement
 - Time measurement for each process

```
double start, exectime;  
MPI_Barrier(MPI_COMM_WORLD);  
start = MPI_Wtime();  
.....  
exectime= MPI_Wtime() - start;
```

In main function

- Execution time of program: the one of the slowest process

FIN CM2

64

65

65